

Parallelized maze routing

王領崧

Outline


- Background
- Problem Description
- Implementation
- Evaluation

Outline

- Background
- Problem Description
- Implementation
- Evaluation

What is maze routing ?

- Commonly used in global routing or detailed routing in EDA tool
- Briefly speaking, find the minimum cost between pins in grid graph
- Easy to implement but time-consuming



This project

- Using CUDA and openMP to accelerate maze routing.
- Implement the idea from the paper “GAMER: GPU Accelerated Maze Routing”

Outline

- Background
- Problem Description
- Implementation
- Evaluation

Problem Description

- On a given weighted graph, given k pins, find a shortest path connecting k pins.
- Input
 - H: grid graph height
 - W: grid graph width
 - Pins: pin's (x, y) coordinate
 - Weights: vertical / horizontal
- Output : A k-pin shortest path
 - Path: all (x, y) coordinate


```
3 3
2
Pin 1 1
Pin 0 2
Vertical 24 26
Vertical 47 13
Vertical 70 2
Horizontal 33 68 51
Horizontal 80 24 77
```

```
1 0 2
2 1 1
3 1 2
```

Problem Description

```
3 3
2
Pin 1 1
Pin 0 2
Vertical 24 26
Vertical 47 13
Vertical 70 2
Horizontal 33 68 51
Horizontal 80 24 77
```

```
1 0 2
2 1 1
3 1 2
```




Outline

- Background
- Problem Description
- Implementation
- Evaluation

Algorithm

1. Choose one pin as source and set its cost to 0
2. Relaxing the cost of other grids using sweep operations
3. Find out the unrouted pin with minimum cost
4. Retrace its path and set the cost to 0
5. Iteratively performing step 2-4 until all the pins are routed



Sweep

- Dynamic programming
- Horizontal / Vertical with two directions

Algorithm 1: $O(n)$ Dynamic Programming Sweep

input : distance d , wire cost c

output: updated distance d

1: **for** $i \leftarrow 1$ **to** $n - 1$ **do**

2: $d(i) \leftarrow \min\{d(i), d(i - 1) + c(i)\}$

3: **end for**


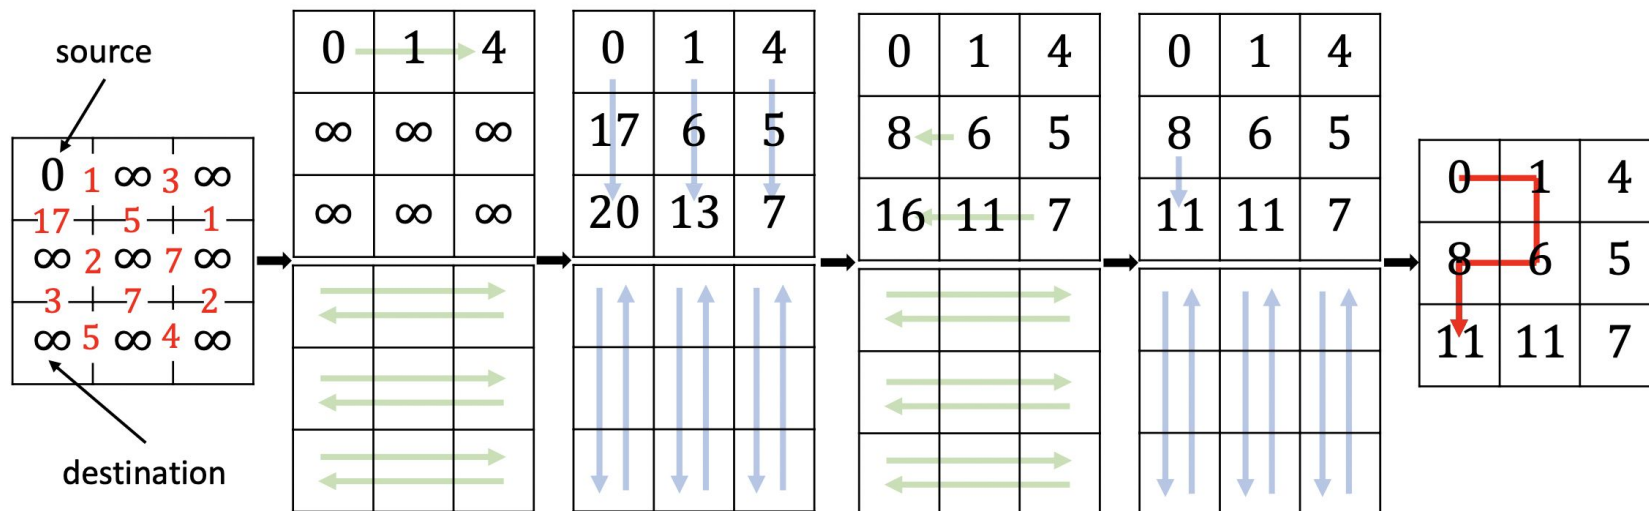


Fig. 3: Sweep

Sweep

- Iteratively performing vertical / horizontal sweeps with direction change.
- Until the grid cost are all minimal.



Sweep with divide and conquer (completed)

- Different rows and columns can work in parallel
- Parallelize inside the row and column (Groups are independent as well)

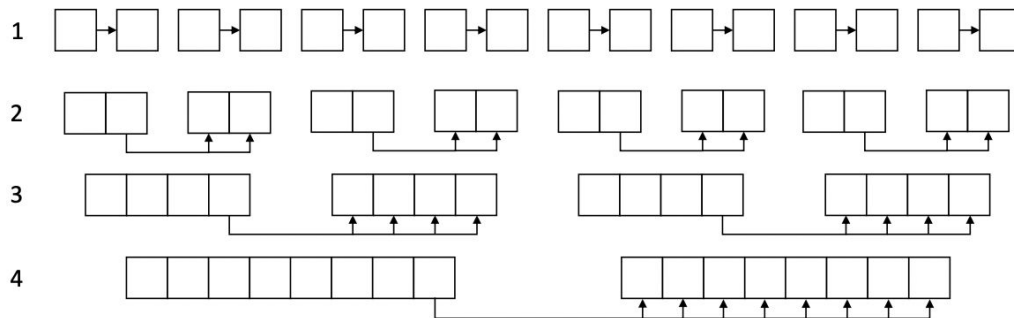


Fig. 5: Parallelizable Sweep

Sweep with reformulation (not completed)

- $d(i)$: cost, $c(i)$: weight
- Turn into two prefix problems

$$s(i) = \sum_{j=0}^i c(j)$$

$$d(i) = \min_{0 \leq j \leq i} \left(d(j) + \sum_{k=j+1}^i c(k) \right) \quad (1)$$

$$d(i) - s(i) = \min_{0 \leq j \leq i} (d(j) - s(j)) \quad (2)$$

Sweep with reformulation

- $d(i)$: cost, $c(i)$: weight, $s(i)$: prefix sum weight
- Prefix sum can be calculated by parallel exclusive

1	2	3	4	5	6	7	8
1	3	3	7	5	11	7	15
1	3	3	10	5	11	7	26
1	3	3	10	5	11	7	36

(a) Phase 1

1	3	3	10	5	11	7	0
1	3	3	0	5	11	7	10
1	0	3	3	5	10	7	21
0	1	3	6	10	15	21	28

(b) Phase 2

$$d(i) - s(i) = \min_{0 \leq j \leq i} (d(j) - s(j)) \quad (2)$$

$c(i)$: wire cost between G-cell $i - 1$ and G-cell i

$d(i)$: distance to G-cell i

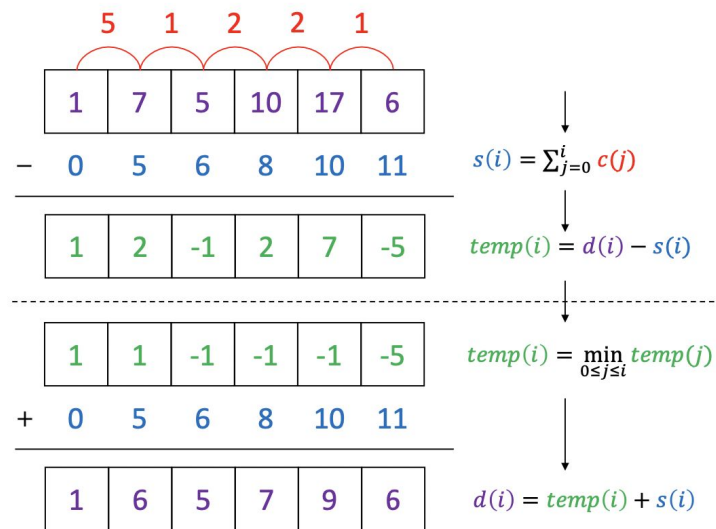
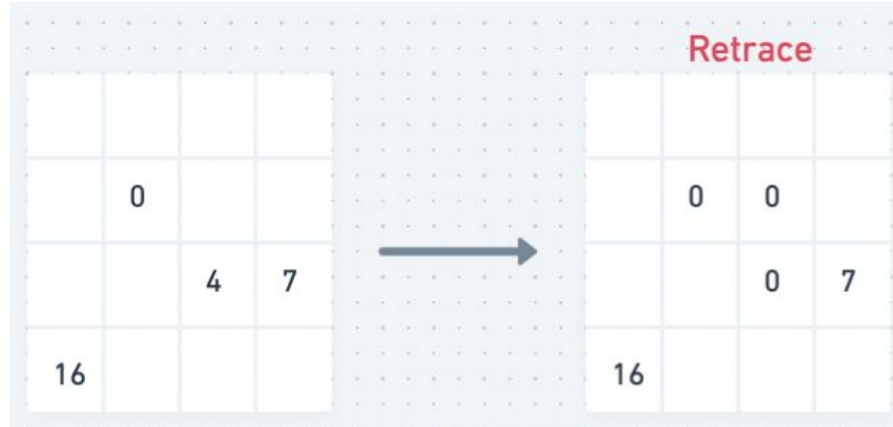


Fig. 6: Sweep by Prefix Sum and Prefix Min

Algorithm

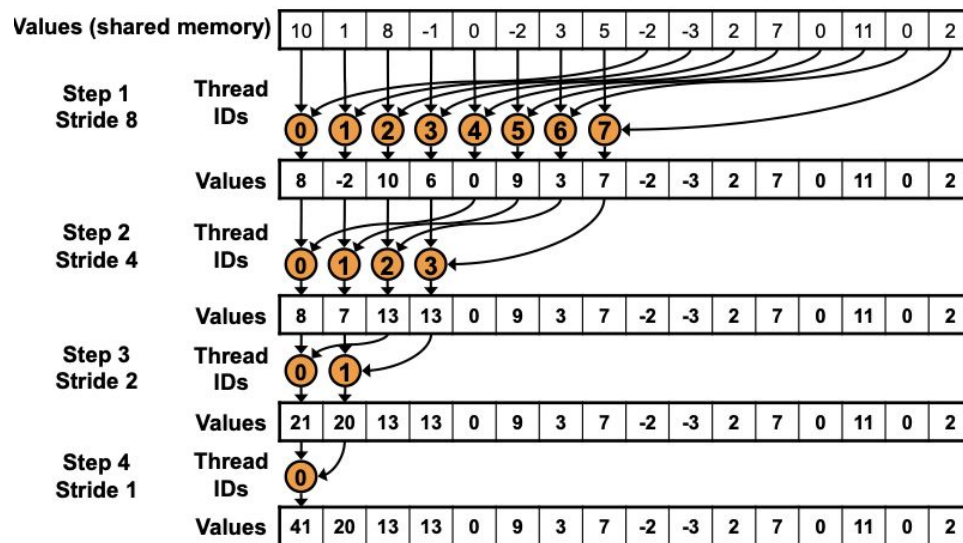
1. Choose one pin as source and set its cost to 0
2. Relaxing the cost of other grids using sweep operations
3. Find out the unrouted pin with minimum cost
4. Retrace its path and set the cost to 0
5. Iteratively performing step 2-4 until all the pins are routed



Get minimum cost pin

- Reduction

Parallel Reduction: Sequential Addressing



Sequential addressing is conflict free

14

Outline


- Background
- Problem Description
- Implementation
- Evaluation

Evaluation

- Omp: Apollo (6 CPUs)
- Cuda: Hades (1 GPU)
- Omp speedup: 3x
- Cuda speedup: 36x


Testcase	W	H	Pins	bfs	sweep	omp	cuda	omp speedup	cuda speedup
1	256	256	4	0.0945	0.3579	0.1331	0.0289	2.69	12.38
2	256	256	8	0.2104	0.6984	0.2658	0.057	2.63	12.25
3	256	256	16	0.3708	1.4031	0.5179	0.1176	2.71	11.93
4	512	512	4	0.5953	6.0169	1.5856	0.2339	3.79	25.72
5	512	512	8	1.3021	11.2363	2.9808	0.3885	3.77	28.92
6	512	512	16	2.6955	24.3993	6.4435	0.8476	3.79	28.79
7	1024	1024	4	6.629	70.5602	20.0804	1.9302	3.51	36.56
8	1024	1024	8	11.0103	135.697	37.5864	3.7695	3.61	36
9	1024	1024	16	22.2053	292.951	79.7602	8.0401	3.67	36.44

Why is sweep much slower than bfs ?



- In simple 2D routing, sweep iterations will be too much
- In practical, via cost (green lines) will reduce the iterations significantly.

Testcase	W	H	Pins	bfs	sweep	omp	cuda
1	256	256	4	0.0945	0.3579	0.1331	0.0289
2	256	256	8	0.2104	0.6984	0.2658	0.057
3	256	256	16	0.3708	1.4031	0.5179	0.1176
4	512	512	4	0.5953	6.0169	1.5856	0.2339
5	512	512	8	1.3021	11.2363	2.9808	0.3885
6	512	512	16	2.6955	24.3993	6.4435	0.8476
7	1024	1024	4	6.629	70.5602	20.0804	1.9302
8	1024	1024	8	11.0103	135.697	37.5864	3.7695
9	1024	1024	16	22.2053	292.951	79.7602	8.0401



Summary

- Complete sweep divide & conquer version
 - speedup: 36x